



PDF Download
3636243.3636249.pdf
28 January 2026
Total Citations: 45
Total Downloads: 2290

Latest updates: <https://dl.acm.org/doi/10.1145/3636243.3636249>

RESEARCH-ARTICLE

Patterns of Student Help-Seeking When Using a Large Language Model-Powered Programming Assistant

BRAD E SHEESE, Illinois Wesleyan University, Bloomington, IL, United States

MARK H. LIFFITON, Illinois Wesleyan University, Bloomington, IL, United States

JAROMÍR ŠAVELKA, Carnegie Mellon University, Pittsburgh, PA, United States

PAUL DENNY, University of Auckland, Auckland, AUK, New Zealand

Open Access Support provided by:

Illinois Wesleyan University

University of Auckland

Carnegie Mellon University

Published: 29 January 2024

[Citation in BibTeX format](#)

ACE 2024: Australian Computing
Education Conference
January 29 - February 2, 2024
NSW, Sydney, Australia

Patterns of Student Help-Seeking When Using a Large Language Model-Powered Programming Assistant

Brad Sheese

bsheese@iwu.edu

Illinois Wesleyan University
Bloomington, Illinois, USA

Jaromir Savelka

jsavelka@cs.cmu.edu

Carnegie Mellon University
Pittsburgh, Pennsylvania, USA

Mark Liffiton

mliffito@iwu.edu

Illinois Wesleyan University
Bloomington, Illinois, USA

Paul Denny

paul@cs.auckland.ac.nz

The University of Auckland
Auckland, New Zealand

ABSTRACT

Providing personalized assistance at scale is a long-standing challenge for computing educators, but a new generation of tools powered by large language models (LLMs) offers immense promise. Such tools can, in theory, provide on-demand help in large class settings and be configured with appropriate guardrails to prevent misuse and mitigate common concerns around learner over-reliance. However, the deployment of LLM-powered tools in authentic classroom settings is still rare, and very little is currently known about how students will use them in practice and what type of help they will seek. To address this, we examine students' use of an innovative LLM-powered tool that provides on-demand programming assistance without revealing solutions directly. We deployed the tool for 12 weeks in an introductory computer and data science course ($n = 52$), collecting more than 2,500 queries submitted by students throughout the term. We manually categorized all student queries based on the type of assistance sought, and we automatically analyzed several additional query characteristics. We found that most queries requested immediate help with programming assignments, whereas fewer requests asked for help on related concepts or for deepening conceptual understanding. Furthermore, students often provided minimal information to the tool, suggesting this is an area in which targeted instruction would be beneficial. We also found that students who achieved more success in the course tended to have used the tool more frequently overall. Lessons from this research can be leveraged by programming educators and institutions who plan to augment their teaching with emerging LLM-powered tools.

CCS CONCEPTS

• **Social and professional topics** → **Computer science education**; **Software engineering education**; • **Human-centered computing** → **Interactive systems and tools**.



This work is licensed under a [Creative Commons Attribution International 4.0 License](https://creativecommons.org/licenses/by/4.0/).

ACE 2024, January 29–February 02, 2024, Sydney, NSW, Australia
© 2024 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-1619-5/24/01.
<https://doi.org/10.1145/3636243.3636249>

KEYWORDS

Intelligent tutoring systems, Intelligent programming tutors, Programming assistance, Novice programmers, Natural language interfaces, Large language models, Guardrails

ACM Reference Format:

Brad Sheese, Mark Liffiton, Jaromir Savelka, and Paul Denny. 2024. Patterns of Student Help-Seeking When Using a Large Language Model-Powered Programming Assistant. In *Australian Computing Education Conference (ACE 2024)*, January 29–February 02, 2024, Sydney, NSW, Australia. ACM, New York, NY, USA, 9 pages. <https://doi.org/10.1145/3636243.3636249>

1 INTRODUCTION

Growing enrollments in introductory programming courses present a number of pedagogical challenges [25]. In particular, providing on-the-spot assistance when students need help becomes a difficult task as student-to-instructor ratios increase. Moreover, not all students feel equally comfortable seeking in-person help from an instructor, and prior work has shown that online options can be more effective for students with low confidence [13]. Thus, there is great interest in finding scalable approaches for providing on-demand, adaptive guidance appropriate for diverse cohorts learning computing [16].

Large language models (LLMs) have emerged quite recently and have shown great potential for generating high quality, human-like feedback. Research in computing education has shown that they can be used to generate programming exercises, natural language explanation of code and more understandable error messages [21]. While these findings suggest promising future applications for real-time student support, concerns have been raised about students over-relying on LLM-based tools, especially using them to generate solutions directly. Indeed, prior work has explicitly called for the use of “guardrails” to be added to educational environments in order to prevent inappropriate use by students [10]. When thinking about the deployment of an LLM-based teaching assistant in the classroom, several critical questions come to mind: How do students interact with the digital assistant in real-world educational settings? Do students predominantly ask it for help to rectify coding errors, or do they also seek deeper conceptual insights? What types of questions dominate, and what information do students generally provide to facilitate the tool’s responses? Is there any evidence that using such a tool helps or harms student learning?

Finding answers to these questions has important educational implications. First, by identifying the kinds of areas in which students seek the most help, instructors can better tailor their teaching methodologies and interventions. The nature of queries can also shed light on potential gaps in the curriculum or teaching materials. Second, the clarity with which students frame their queries can provide insights into their metacognitive skills as well as their abilities to clearly express issues they are encountering. Especially now that communication with AI-powered tools is becoming commonplace, the ability to craft prompts that are clear and that contain sufficient information is becoming an essential skill [7]. Last, evaluating the relationship between consistent tool engagement and course performance can offer insights into the tangible benefits of integrating LLM-powered tools into the classroom.

In this paper we investigate student use of an LLM-powered teaching assistant (TA) we developed to provide real-time help for students learning programming. We deployed the tool, CodeHelp, in an introductory computer and data science course throughout the Spring 2023 semester [22]. One of the key contributions of CodeHelp is its effective use of robust “guardrails” that are specifically designed to prevent it from directly outputting solution code. Thus, the tool offers students a means of assistance while attempting to avoid the over-reliance that direct use of LLMs may promote.

During the semester, students were free to make their own choices about when and how to use CodeHelp, although when it was introduced in the course we did encourage students to seek help from CodeHelp first, before asking the instructor or TAs. Our aim with this research is to understand how students make use of an always-available LLM-powered TA over a semester-long course. Our evaluation is guided by the following research questions:

- (RQ1) What kind of help do students seek when using CodeHelp, and how do they construct their queries?
- (RQ2) To what extent does student usage of CodeHelp over a semester-long introductory course correlate with their final course performance?

Interest in LLMs and how they can be productively used in education is rapidly growing. To date, although there are reports of new and existing tools that have integrated LLMs, these have not been deployed at scale over an entire semester in a computing classroom. With this current paper, we make two novel contributions to the computing education literature. We:

- (C1) analyze how students use an always available LLM-powered tutor during an authentic semester-long course.
- (C2) investigate how student usage of the LLM-powered tutor is associated with their performance on the course.

2 RELATED WORK

Large Language Models (LLMs) have been applied to a wide range of tasks relevant to computing education, including solving code exercises [7, 12, 35], answering [32, 33] and generating programming MCQs [11], detecting bugs [23], writing tests [4, 17], and generating code explanations [20, 24] and programming exercises [31]. The breadth of this work demonstrates both the flexibility and the capabilities of LLMs. However, despite their impressive performance at many tasks, LLMs may not be as effective as human tutors in some domains [26, 32]. For example, a study by Pardos

and Bhandari that compared the efficacy of hints generated by ChatGPT and by human tutors in the subject of algebra found that significant learning gains were only observed for students who were provided human-created hints [28]. Thus, there is an important need to carefully study the use of LLM-powered tutors in the context of computing classrooms.

Prior research on the use of AI tools in programming education has raised concerns around student over-reliance on AI-generated code [2, 3, 6, 30]. For example, Kazemitabaar et al. examined student usage of the Coding Steps tool, which blends AI code generation with an online platform, in a supervised lab environment [18]. They found that when code was generated by the AI model, around half the time it was submitted by students without them making any changes to it. There are also concerns about the accuracy of the responses generated by AI tools and their suitability for providing unmoderated assistance to students. Hellas et al. studied responses generated by Codex and GPT-3.5 to 150 student help requests from an existing dataset [15]. The original data came from a platform where requests for help from students were manually responded to by a teacher. Their results showed that neither of the AI models found all of the issues in the students’ help requests and, concerning, false positives were common. In very similar work, Balse et al. found high variability in the accuracy of LLM-generated feedback on student submissions [1].

Our current work extends prior work in several important ways. Firstly, we have designed guardrails in CodeHelp to prevent it from revealing code solutions directly, thus avoiding the undesirable usage patterns observed by Kazemitabaar et al. in which students simply submitted the AI-generated code [18]. Secondly, we conduct our evaluation over 12 weeks and explore how students interact with it outside of scheduled class sessions rather than in narrower time-spans and contexts. Thirdly, we analyse the help requests created by students as opposed to the responses provided by the LLM, as this is already a widely studied topic [5, 12, 32].

3 CODEHELP

We created CodeHelp to be an automated assistant for students in programming and CS courses [22]. It responds to semi-structured student queries using a large language model with guardrails; specifically, it will not provide solution code in its responses regardless of how a student writes their query. It is designed to be an always-and instantly-accessible source of support similar to what a TA might provide, intended to augment and complement the support provided by course instructors and TAs.

For broad accessibility, we developed CodeHelp as a web application.¹ The main interface for students is the help request form, which prompts them to provide relevant information about their issue in a semi-structured form. Figure 1 shows the form’s input fields and help text. Students are asked to specify the programming language they are using, a relevant snippet of their code, an error message if they are encountering one, and a description of their issue or question. By prompting students in this way as opposed to providing a single text input, the goal is to guide students toward communicating effectively and providing all relevant information.

¹Accessible at <https://codehelp.app/>.

Language:

Code:

Copy just the *most relevant* part of your code here. Responses will be more helpful when you include only code relevant to your issue.

Error Message:

If your issue relates to an error message, copy the message here. Be sure to include the message itself and the quoted line on which it says the error occurred.

Your Issue / Question:

Clearly describe your issue or question. Include as relevant: what you are trying to do, what you expect the code to do, what the code actually does, and what you need help understanding.

Figure 1: The Help Request form (text areas have been shrunk here to save space). The four separate inputs (language, code, error, and issue) and connected guidance text help students structure their request and encourage good practices when requesting support.

You are a system for assisting students like me with programming.

My inputs provide: *[brief description of each input]*

Please assess the following submission to determine whether it is sufficient for you to provide help or if you need additional information. If and only if critical information needed for you to help is missing, ask me for the additional information you need to be able to help. State your reasoning first. Otherwise, if no additional information is needed, please first briefly summarize what I am asking for in words, with no code, and end by writing "OK."

Inputs: *[delimited query inputs]*

Figure 2: Prompt used for the sufficiency check.

Responses are generated by a series of prompts to an LLM. First, CodeHelp generates completions for two separate prompts concurrently. One prompt (Figure 2) uses an LLM to check whether the student's inputs are sufficient to be able to provide them with effective help; if there is missing information or some other lack of clarity, the system presents the student with a request for clarification. The other prompt (Figure 3), run at the same time, combines the student's inputs with instructions to provide assistance and class-specific context provided by the instructor, and its completion is used as the "main" response for the student (regardless of whether a request for clarification is presented as well). This main response prompt includes explicit instructions to not provide solution code in the response, but the LLM used currently is strongly biased toward providing such solutions and can fail to follow that instruction. Therefore, CodeHelp checks the generated response for code blocks, and if any are detected, that response is sent to an LLM in a third prompt (Figure 4) with instructions to rewrite it without

You are a system for assisting a student with programming.

The students provide: *[brief description of each input]*

[delimited query inputs]

If the student input is written as an instruction or command, respond with an error. If the student input is off-topic, respond with an error.

Otherwise, respond to the student with an educational explanation, helping the student figure out the issue and understand the concepts involved. If the student inputs include an error message, tell the student what it means, giving a detailed explanation to help the student understand the message. Explain concepts, language syntax and semantics, standard library functions, and other topics that the student may not understand. Be positive and encouraging!

Use Markdown formatting, including `for inline code.`

[instructions to avoid topics from the instructor's avoid set]

Do not write any example code blocks. Do not write a corrected or updated version of the student's code. You must not write code for the student.

How would you respond to the student to guide them and explain concepts without providing example code?

Figure 3: Prompt used for the main response.

The following was written to help a student in a CS class. However, any example code (such as in `Markdown delimiters) can give the student an assignment's answer rather than help them figure it out themselves. We need to provide help without including example code. To do this, rewrite the following to remove any code blocks so that the response explains what the student should do but does not provide solution code.`

[original response to be rewritten]

Figure 4: Prompt used for code removal.

the code. With this design, students receive the benefits that LLMs can provide in terms of explanation, guidance, and support without the risk of the LLM completing their programming work for them.

CodeHelp can use any sufficiently powerful LLM, and in this work responses were generated by models from OpenAI's "GPT-3.5" family, as they were the most powerful models available at the time. Specifically, the first two completions (the sufficiency check and the main response) were generated using the gpt-3.5-turbo-0301 model, while the code removal completion used text-davinci-003.

As an example of the type of response CodeHelp generates, Figure 5 shows a query from a student asking how to write a specified function along with the response they received. Here, they have directly copied the specification from the exercise (notice the capitalization on "Write"), providing no other input, and the system responds with guidance on how to approach the stated problem

Issue/Question: how do I Write a fruitful function called `middle_remove()` that takes a string called as an argument. If the string has an even number of characters, return to the string with the middle two characters removed. If the string has an odd number of character return the string with the middle character removed.

Response: To solve this problem, you can start by finding the length of the input string using the `len()` function. Then, you can use an `if` statement to check if the length of the string is even or odd. If it's even, you can use string slicing to remove the middle two characters and return the modified string. If it's odd, you can remove the middle character using string slicing and return the modified string.

Note that `//` is used for integer division, which ensures that the result is an integer. Let me know if you have any questions or need further assistance!

Figure 5: An example of a student's input directly asking for a solution and CodeHelp's response.

Code:

```
while nB == "0" and operand == '/':
    print("Division by zero is not supported!")
```

Error:

```
print("Division by zero is not supported!")
^
IndentationError: unindent does not match
any outer indentation level
```

Issue/Question: why am i getting this error

Response: This error occurs because the `print()` statement is not indented correctly. In Python, indentation is used to denote code blocks, and the code block that the `print()` statement is in is not indented to the same level as the `while` loop. To fix this, make sure that the `print()` statement is indented one level further than the `while` loop.

Figure 6: An example of a student's input asking for help with an error and CodeHelp's response.

but does not write any code. Figure 6 shows another example, in which a student provides some code and an error message they have received, trying to understand it. These generally exemplify how CodeHelp responds to queries: it provides guidance for solving problems, including high-level algorithmic steps, and it explains concepts as needed.

4 METHODS

We evaluated CodeHelp in two sections of an undergraduate introductory computer- and data-science course ($n = 52$) taught by an author of this paper in the Spring semester of 2023. The course

serves students from across the institution who take the course to meet general education requirements or to meet requirements for data-analytic or data-science related credentials. The course is taught in a “flipped” style with students spending much of the class time working through assignments on lab computers. The instructor and a TA assist students and provide instruction and support as needed. CodeHelp was introduced in the fourth week of the semester with a quick demonstration in class. During class, students were encouraged to use CodeHelp for assistance first before asking the instructor or TA for help, but they were otherwise free to make their own choices about when and how to use it.

After the course concluded, we used the complete set of student query logs to perform several analyses related to our research questions. Our analyses considered 49 students who submitted at least one query during the course of the semester. Students submitted a total of 2,591 queries.

4.1 CodeHelp Overall Usage

To capture different aspects of students' CodeHelp usage a composite measure was created from the three metrics described below. Each metric was transformed to help address skew and then standardized before being averaged together. The composite measured showed good internal consistency (Cronbach's $\alpha = .87$).

- (1) Total Queries: We calculated the total number of queries a student submitted during the course of the semester. While, on average, students submitted about 40 queries during the semester ($M = 39.50$, $SD = 42.45$), one student was an extreme outlier submitting a total of 614 queries. This student was excluded from subsequent correlational analyses examining overall usage.
- (2) Total Sessions: This represents the unique number of ‘sessions’ in which a student engaged with CodeHelp during the semester ($M = 13.50$, $SD = 11.54$). A ‘session’ is characterized as a span of activity that is distinctly segregated by a minimum inactivity period of one hour.
- (3) Average Length of Session: We calculated the mean length of a student's sessions during the semester in seconds ($M = 885.13$, $SD = 1082.17$), which was roughly 15 minutes.

4.2 Characterizing Queries

To investigate how students used CodeHelp, we characterized their queries using a mix of manual categorization and automated analysis.

4.2.1 Deduplication and Query Cleaning. Students would sometimes resubmit a query with little or no change, likely in an attempt to receive a different response from the LLM, which is non-deterministic. We filtered out these duplicate queries prior to content coding so that repetitions did not exert an undue influence on efforts to characterize the results². Of the 2,591 raw queries, this process filtered 509 as duplicates, leaving 2,082 queries to be coded. In addition, during the coding process, three queries were identified

²Any two consecutive queries X and Y from a given user were considered duplicate based on a measure of similarity between X and Y . Similarity was calculated by computing the Levenshtein edit distances between each of the code, error, and issue strings of X and Y , normalize each to the range $[0, 1]$ by dividing by the longer of the two strings in each case, and add the three normalized distances. If the sum is less

that did not pertain to the course (e.g. “What is the meaning of life?”) and were omitted from subsequent analysis.

4.2.2 Query Content Categories. We manually coded all 2,082 duplicated queries into categories based on their contents, with each query coded independently by two of the authors. To determine the categories, all four authors independently read different selections of the queries, each creating categories as needed to classify successive queries until reaching saturation (a long stretch of queries that required no new categories). In total, this effort covered over 500 different queries. Then, the authors compared and refined their proposed taxonomies. This process produced four major categories of queries:

- (1) “Debugging”: Queries where students were looking for help to solve specific errors and faults in their code. Debugging queries were further sub-categorized into one of the following types: a) debugging queries that only included the error b) debugging queries that only included the desired outcome (what the code is supposed to do), or c) debugging queries that included both the error and the desired outcome.
- (2) “Implementation”: Queries about implementing code or functions to solve specific assignment problems. These queries included code and/or referenced course assignment instructions.
- (3) “Understanding”: Queries centered around gaining an understanding of programming concepts, algorithms, data structures, language, or library features, but not obviously asking how to complete a given assignment problem. These queries did not include code or reference assignment instructions.
- (4) “Nothing” (no explicit student request): Queries that provided no error or meaningful issue. These queries commonly only provided code with no other context.

4.2.3 Inter-Rater Reliability. To examine inter-rater reliability of human coded categories, we calculated Cohen’s κ for all categories and sub-categories, for all categories with the sub-categories collapsed, and then for individual categories using binary comparisons. Overall reliability for ratings of all categories and sub-categories was substantial ($\kappa = .75$), and overall reliability was even higher ($\kappa = .83$) when Debugging sub-categories were collapsed in a single Debugging category [27]. κ values for individual category ratings are presented in Table 1.

4.2.4 Identifying “Low-Effort” Queries. During the course and while inspecting queries later, we noticed that students would often submit “low-effort” queries in which they wrote little or nothing themselves. Some queries contained little or nothing in the Issue input at all, while others were mostly or entirely copied and pasted from the instructions in a class exercise or assignment. This behavior somewhat mirrors the low-effort submissions noted by Kazemitabaar et al. in their LLM-powered tool [18]. To investigate this further, we

than a threshold k , we considered the second query to be a duplicate of the first.

$$\text{Duplicate}(X, Y) := \left(\sum_{s \in \{\text{code}, \text{error}, \text{issue}\}} \frac{\text{Levenshtein}(s_X, s_Y)}{\max(|s_X|, |s_Y|)} \right) < k$$

The sum can range from 0 (all three strings equivalent) to 3 (all three strings entirely different). We chose a threshold k of 0.25, meaning, roughly, that a query is considered a duplicate of a student’s previous query if it differs in no more than a quarter of one of the query’s three strings.

used automated analyses to identify queries with two additional characteristics:

(1) *Empty or Short Issue*

Any issue text whose length was less than 10 characters was classified as empty/short, matching 8% of all queries. In these cases, there is little to indicate the student’s intent, but CodeHelp’s responses typically described what the code appeared to be aiming to do and pointed out potential errors in the code. Students may have been submitting these queries expecting that type of response. Note that these were inferred to be Debugging queries in the case that something was specified in the error input.

(2) *Substantially Copied from Course Materials*

To quantify copied course content, we computed “diffs” between each query’s Issue text and the text of every exercise in the course using Python’s `difflib` library. For each diff, we calculated the percentage of the Issue text that matched the exercise, and we took the highest percentage across all exercises as the “copied percentage” for that Issue text. We chose a threshold of 80% as a cutoff to indicate that most if not all of the Issue text was directly copied from an exercise. In cases where the match percentage was less than 100% but greater than 80%, the non-matching text was often something like “How do I” followed by the copied exercise text or a simple bit of added context like “This is using Pandas.” 14% of all queries were above this 80%-matching threshold.

4.3 Course Performance

To address the question of how CodeHelp usage relates to performance in the class we used course points that were earned throughout the semester through typical course activities including quizzes, assignments, and reading responses. Following transformation to correct skew, the course points earned in each activity were z-scored and then aggregated to create an overall Course Performance metric.

5 RESULTS

5.1 Characterizing Queries


Figure 7 presents the fraction of each student’s queries (after deduplication) in each category, including both the manually-coded and the automatically-assigned categories.

Students tended to submit more debugging queries than any other type. On average, 48% of a student’s queries were debugging queries, and 47 of the 49 students submitted at least one debugging query. For 21 students, debugging made up more than half of all of their queries. We do see a wide range here, though, and in fact students are evenly distributed across the entire range from 0% to 100% in that category. Of the debugging subcategories, including the error but not expected outcome was the most common, making up 58% of all debugging queries. The remaining 42% of debugging queries contained or at least suggested an expected outcome, while only 31% of debugging queries contained both an indication of the erroneous behavior and the expected outcome.

Implementation queries were broadly popular, though less so than debugging. On average, 38% of a student’s queries were requesting implementation help. 44 of 49 students submitted at least

Table 1: Query Categories. Descriptives (count and %) and inter-rater agreement for top-level content categories and debugging subcategories.

Query Category	Example Text	Count	Percent	Kappa
Debugging (all)	“Why doesn’t/can’t/isn’t the code doing X?”	833	40%	.86
Implementation	“How do I get it to X?”	1038	50%	.83
Understanding	“What is the difference between functions X and Y?”	161	8%	.77
Nothing	(No explicit student request)	47	2%	.78
Debugging Sub-Categories		Count	Percent	Kappa
Including error		484	58%	.78
Including outcome		90	11%	.41
Including error & outcome		259	31%	.65

**Figure 7: Fraction of each student’s queries (after deduplication) with each characteristic. Each point represents one student and the fraction of their queries having that characteristic. The box-and-whiskers plots follow convention: the left and right edges of each box indicate the lower and upper quartiles, the red line inside the box indicates the median, and whiskers are drawn to the farthest point that is at most 1.5 times the inter-quartile range beyond the box.**

one implementation query, and for 18 students, more than half of their queries were implementation queries.

Queries related to understanding were relatively rare, with only 28 students submitting at least one query of this type. On average, just 9% of each student’s queries were understanding related. For two students, understanding queries were the majority of their queries, though both of these students submitted fewer than ten queries, total.

In many cases, students submitted a query with nothing or very little (fewer than ten characters) in the Issue input. 29 students did this at least once, and for five students, the Issue input was effectively empty in more than 10% of their queries. Similarly, about 14% of all queries were entirely or substantially copied from exercise instructions (at or above the 80% threshold), and 31 of the

49 students submitted at least one of these queries. Six students submitted substantially copied Issue text in more than 50% of their queries. Taking these two categories together as queries in which the student wrote little or nothing themselves in the Issue input, we found that 41 students submitted at least one query with either of those characteristics, and such queries made up over 50% of submitted queries for 12 students.

All together, these data illustrate that, broadly, the students in this course tended to be:

- Primarily focused on receiving immediate help with a programming problem rather than on improving their understanding of concepts – indicated by far higher incidence of Implementation and Debugging queries compared to Understanding.
- Unsophisticated in their requests, often writing little or nothing themselves to direct the assistance they receive, and often omitting important details (any indication of either the error or the expected outcome) when requesting debugging help.

5.2 Tool Usage and Course Performance

We computed Pearson’s r correlation coefficient to analyze the relationship between tool usage and course performance ($r = 0.38$, $p = 0.0126$). This result indicates a modest positive relation between the two variables. As illustrated in Figure 8, students showing higher tool usage tended to show somewhat higher course performance. Please note that this correlation does not suggest, or support, any claims of causality and should not be interpreted as support for the claim that usage caused higher performance.

6 DISCUSSION & CONCLUSIONS

This work examined how students seek help when using an LLM-powered teaching assistant during a semester long introductory programming course. We found that individual students showed a variety of approaches when constructing queries, but that the majority of queries concerned debugging code errors or getting started with programming assignments. Students were relatively unlikely to request help with deepening their understanding of general concepts or principles related to their assignments. We found the scope and quality of student queries to be limited in




Figure 8: Scatter plot of the association between CodeHelp Usage and Course Performance. See separate Usage and Course Performance sections for details on metric transformation, standardization, and aggregation. **Note:** Due to standardization, both metrics have a mean of 0 and a standard deviation of 1.

many cases. For example, we found many instances of debugging queries without descriptions of problems or intended behavior.

Our results highlight the strengths and weaknesses of these types of tools in supporting student learning and problem-solving. Students collectively submitted thousands of queries throughout the course of the semester without any extrinsic incentive to do so. Many of the queries were well-articulated and provided good opportunities for the student to learn from the response. However, the poor quality of many queries also highlights the need to cultivate effective help-seeking skills. Guardrails that prohibit LLMs from providing code may be just the first step to guiding students through the process of problem-solving and fostering an inquisitive mindset rather than mere solution-seeking.

We also examined how student performance in the course was related to their overall pattern of CodeHelp usage during the semester. We were particularly interested in this relationship given the widely expressed concerns in many areas of education that the use of LLMs will undermine student learning [29]. Indeed, recent work in computing education has clearly illustrated the risks in providing students with unconstrained access to LLMs. For example, Prather et al. observed students using Github Copilot to solve a typical CS1 assignment in a laboratory setting [30]. Interviews with participants were revealing, with one-third of students expressing concerns about the potential negative effects of LLM-based coding tools on their learning. These included concerns around solving problems without needing to understand or even process what is going on, and worries that using such tools would harm problem solving skills given that the tools “would just do it for me”. Similarly, Kazemitabaar et al. performed a thematic analysis of student queries to an AI code generator based on OpenAI’s Codex

model [19]. They identified four distinct approaches that students used, and found that by far the most common approach (AI Single Prompt) involved learners prompting Codex just once to generate the entire solution to a task. Not only was this the most commonly used approach, but it also correlated negatively with student performance on subsequent post-tests, suggesting it was detrimental to learning.

In light of these well-documented concerns, our goal with CodeHelp was to provide the benefits of LLM assistance to students while reducing the harms (e.g., over-reliance). In the current study, we found no evidence consistent with the idea that CodeHelp usage undermined student learning. Instead, our results showed a modest positive association between overall use during the semester and course performance. Further empirical work, using larger sample sizes and a diversity of course designs, will be required before any firm conclusions can be drawn. However, these preliminary findings suggest that when appropriately integrated into the educational framework, tools like CodeHelp may enhance the learning experience without diminishing the student’s ability to learn and apply new concepts independently.

An interesting avenue for future work would be to provide automatic feedback to students directly within the tool to help them construct effective queries when seeking help. This aligns with the goals of recent work by Denny et al., on the use of Prompt Problems in introductory programming courses, which leverages LLMs to teach students how to craft good prompts for AI-based code generators [8, 9]. In the case of our work with CodeHelp, if a student failed to clearly describe the intended behaviour of the code for which they are seeking help, the tool could prompt them to do so before providing assistance. Similar feedback could be used to dissuade students from relying too much on the tool for certain kinds of queries, such as ‘implementation’ questions. In some cases, such as where an input field is left empty, generating suitable feedback to students would be straightforward. In other cases, the ability to automatically classify the queries that students submit into the categories we have presented in this work would facilitate the generation of useful feedback. For instance, this would permit the identification of debugging queries where the intended behaviour is missing, or potential overuse of CodeHelp for requesting help with implementation. Recent work by Gao et al. has shown good potential for this type of automated classification [14].

6.1 Limitations and Threats to Validity

Our research context was limited to a specific introductory computer and data science course with a sample of 52 students. Caution is warranted in generalizing our findings to other student demographics, course structures, and institutional cultures. Similarly, the 12-week duration of our investigation may influence both the frequency and type of student queries. It may be the case that the nature of students’ queries changed over time as they gained more experience with the tool. Thus, longer or shorter studies may find other patterns of usage. Similarly, the association we reported between academic performance and usage may be yolked to a variety of factors that we did not address, such as course style, content, difficulty, average student preparedness, and student comfort with AI-driven tools. As this study was conducted in the Spring of 2023,

many students may not have had much, if any, prior experience working with LLMs. As the general population becomes increasingly familiar with LLMs, different patterns of results may emerge.

Our qualitative approach to the development of content categories as well as our use of human raters introduces quite a bit of subjectivity into our study. While we aimed to be objective and to establish the reliability of our approach, the inherent subjectivity cannot be ignored. Our categories and ratings may reflect unknown biases and it is possible that other researchers with different backgrounds might produce different categorizations schemes, different ratings, and, ultimately, different conclusions.

Finally, we acknowledge the role that interface design can play in shaping user behaviors and interactions. For example, CodeHelp’s interface may have steered students towards submitting debugging-focused queries. Overall, our tool’s layout might have been more amenable to code implementation queries, at the expense of more abstract or conceptual inquiries. As we look to future refinements, it is crucial to contemplate redesigning the interface or embedding clear cues to facilitate a broader spectrum of queries, ensuring a balance between technical troubleshooting and conceptual understanding.

6.2 Implications for Teaching Practice and Future Work

Future work should move beyond overall course performance and focus on studying the long-term effects of LLM usage on student problem-solving skills and conceptual development. Furthermore, examining a broader range of students from across computing disciplines and educational levels would provide a more thorough understanding of LLMs’ capabilities in supporting student learning in other contexts.

It is a notable feature of CodeHelp and the LLM-powered approach that even when students do not clearly communicate their issue or even write little to nothing in their own words, they can still receive a helpful response from the tool. However, these helpful responses can then positively reinforce the submission of poorly articulated queries. This is unfortunate, as effective communication of programming problems is a useful skill to develop and is an activity that may help students identify their learning gaps and deepen their conceptual understanding of course materials. While the students appear to prefer immediate task-specific assistance, it is essential to also emphasize the importance of developing deep conceptual understanding. Programming education is not limited to resolving immediate coding challenges, but also entails building a solid foundation of underlying concepts.

There are many opportunities to investigate these tendencies and their implications further. Empirically, there is more to be learned about the balance between the efficiency of “low-effort” queries and the learning costs to the student of not articulating their issue. When further developing tools, additional or alternative methods of scaffolding and guiding students’ inputs may encourage use that helps advance students’ conceptual understanding and communication skills in addition to helping them solve their immediate problems. Automatically detecting the type of query a student is asking, and tailoring the LLM-generated response based on the type has potential to increase the effectiveness of the tool [34].

Additionally, educators should consider offering explicit guidance on formulating high-quality technical questions, e.g., by integrating structured tutorials, examples of well-formulated questions, or feedback mechanisms into the assisting tool. Finally, practicing effective interaction with LLM powered tools as described in [9] may help students overcome these issues.

REFERENCES

- [1] Rishabh Balse, Bharath Valaboju, Shreya Singhal, Jayakrishnan Madathil Warriem, and Prajish Prasad. 2023. Investigating the Potential of GPT-3 in Providing Feedback for Programming Assessments. In *Proceedings of the 2023 Conference on Innovation and Technology in Computer Science Education V. 1* (Turku, Finland) (ITiCSE 2023). Association for Computing Machinery, New York, NY, USA, 292–298. <https://doi.org/10.1145/3587102.3588852>
- [2] Brett A Becker, Paul Denny, James Finnie-Ansley, Andrew Luxton-Reilly, James Prather, and Eddie Antonio Santos. 2023. Programming Is Hard-Or at Least It Used to Be: Educational Opportunities and Challenges of AI Code Generation. In *Proceedings of the 54th ACM Technical Symposium on Computer Science Education V. 1*. ACM, Toronto, ON, Canada, 500–506.
- [3] Peter Brusilovsky, Barbara J Ericson, Cay S Horstmann, and Christian Servin. 2023. The Future of Computing Education Materials.
- [4] Bei Chen, Fengji Zhang, Anh Nguyen, Daoguang Zan, Zeqi Lin, Jian-Guang Lou, and Weizhu Chen. 2022. CodeT5: Code Generation with Generated Tests. <http://arxiv.org/abs/2207.10397> arXiv:2207.10397 [cs].
- [5] Bruno Pereira Cipriano and Pedro Alves. 2023. GPT-3 vs Object Oriented Programming Assignments: An Experience Report. In *Proceedings of the 2023 Conference on Innovation and Technology in Computer Science Education V. 1* (Turku, Finland) (ITiCSE 2023). Association for Computing Machinery, New York, NY, USA, 61–67. <https://doi.org/10.1145/3587102.3588814>
- [6] Jonathan E Collins. 2023. Policy Solutions: Policy questions for ChatGPT and artificial intelligence. *Phi Delta Kappan* 104, 7 (2023), 60–61.
- [7] Paul Denny, Viraj Kumar, and Nasser Giacaman. 2023. Conversing with Copilot: Exploring Prompt Engineering for Solving CS1 Problems Using Natural Language. In *Proceedings of the 54th ACM Technical Symposium on Computer Science Education V. 1*. ACM, Toronto ON Canada, 1136–1142. <https://doi.org/10.1145/3545945.3569823>
- [8] Paul Denny, Juho Leinonen, James Prather, Andrew Luxton-Reilly, Thezyrie Amarouche, Brett A. Becker, and Brent N. Reeves. 2023. Prompt Problems: A New Programming Exercise for the Generative AI Era. arXiv:2311.05943 [cs.HC]
- [9] Paul Denny, Juho Leinonen, James Prather, Andrew Luxton-Reilly, Thezyrie Amarouche, Brett A. Becker, and Brent N. Reeves. 2023. Promptly: Using Prompt Problems to Teach Learners How to Effectively Utilize AI Code Generators. arXiv:2307.16364 [cs.HC]
- [10] Paul Denny, James Prather, Brett A. Becker, James Finnie-Ansley, Arto Hellas, Juho Leinonen, Andrew Luxton-Reilly, Brent N. Reeves, Eddie Antonio Santos, and Sami Sarsa. 2023. Computing Education in the Era of Generative AI. arXiv:2306.02608 [cs.CY]
- [11] Jacob Doughty, Zipiao Wan, Anishka Bompelli, Jubahed Qayum, Taozhi Wang, Juran Zhang, Yujia Zheng, Aidan Doyle, Pragnya Sridhar, Arav Agarwal, et al. 2023. A Comparative Study of AI-Generated (GPT-4) and Human-crafted MCQs in Programming Education. arXiv preprint arXiv:2312.03173 (2023).
- [12] James Finnie-Ansley, Paul Denny, Brett A Becker, Andrew Luxton-Reilly, and James Prather. 2022. The robots are coming: Exploring the implications of openai codex on introductory programming. In *Proceedings of the 24th Australasian Computing Education Conference*. 10–19. <https://doi.org/10.1145/3511861.3511863>
- [13] Zhikai Gao, Sarah Heckman, and Collin Lynch. 2022. Who Uses Office Hours? A Comparison of In-Person and Virtual Office Hours Utilization. In *Proceedings of the 53rd ACM Technical Symposium on Computer Science Education - Volume 1* (Providence, RI, USA) (SIGCSE 2022). Association for Computing Machinery, New York, NY, USA, 300–306. <https://doi.org/10.1145/3478431.3499334>
- [14] Zhikai Gao, Collin Lynch, Sarah Heckman, and Tiffany Barnes. 2021. Automatically Classifying Student Help Requests: A Multi-Year Analysis. *International Educational Data Mining Society* (2021).
- [15] Arto Hellas, Juho Leinonen, Sami Sarsa, Charles Koutchme, Lilja Kujanpää, and Juha Sorva. 2023. Exploring the Responses of Large Language Models to Beginner Programmers’ Help Requests. arXiv preprint arXiv:2306.05715 (2023).
- [16] Yann Hicke, Anmol Agarwal, Qianou Ma, and Paul Denny. 2023. ChaTA: Towards an Intelligent Question-Answer Teaching Assistant using Open-Source LLMs. arXiv:2311.02775 [cs.LG]
- [17] Sajed Jalil, Suzzana Rafi, Thomas D. LaToza, Kevin Moran, and Wing Lam. 2023. ChatGPT and Software Testing Education: Promises & Perils. <http://arxiv.org/abs/2302.03287> arXiv:2302.03287 [cs].
- [18] Majeed Kazemitabaar, Justin Chow, Carl Ka To Ma, Barbara J. Ericson, David Weintrop, and Tovi Grossman. 2023. Studying the Effect of AI Code Generators on Supporting Novice Learners in Introductory Programming. In *Proceedings*

- of the 2023 CHI Conference on Human Factors in Computing Systems (Hamburg, Germany) (CHI '23). Association for Computing Machinery, New York, NY, USA, Article 455, 23 pages. <https://doi.org/10.1145/3544548.3580919>
- [19] Majeed Kazemitabaar, Xinying Hou, Austin Henley, Barbara J. Ericson, David Weintrop, and Tovi Grossman. 2023. How Novices Use LLM-Based Code Generators to Solve CS1 Coding Tasks in a Self-Paced Learning Environment. arXiv:2309.14049 [cs.HC]
- [20] Juho Leinonen, Paul Denny, Stephen MacNeil, Sami Sarsa, Seth Bernstein, Joanne Kim, Andrew Tran, and Arto Hellas. 2023. Comparing Code Explanations Created by Students and Large Language Models. <http://arxiv.org/abs/2304.03938> arXiv:2304.03938 [cs].
- [21] Juho Leinonen, Arto Hellas, Sami Sarsa, Brent Reeves, Paul Denny, James Prather, and Brett A. Becker. 2023. Using large language models to enhance programming error messages. In *Proceedings of the 54th ACM Technical Symposium on Computer Science Education V. 1*. 563–569.
- [22] Mark Liffiton, Brad Sheese, Jaromir Savelka, and Paul Denny. 2023. Codehelp: Using large language models with guardrails for scalable support in programming classes. *arXiv preprint arXiv:2308.06921* (2023).
- [23] Stephen MacNeil, Paul Denny, Andrew Tran, Juho Leinonen, Seth Bernstein, Arto Hellas, Sami Sarsa, and Joanne Kim. 2023. Decoding Logic Errors: A Comparative Study on Bug Detection by Students and Large Language Models. arXiv:2311.16017 [cs.HC]
- [24] Stephen MacNeil, Andrew Tran, Arto Hellas, Joanne Kim, Sami Sarsa, Paul Denny, Seth Bernstein, and Juho Leinonen. 2023. Experiences from Using Code Explanations Generated by Large Language Models in a Web Software Development E-Book. In *Proceedings of the 54th ACM Technical Symposium on Computer Science Education V. 1* (<conf-loc>, <city>Toronto ON</city>, <country>Canada</country>, </conf-loc>) (SIGCSE 2023). Association for Computing Machinery, New York, NY, USA, 931–937. <https://doi.org/10.1145/3545945.3569785>
- [25] National Academies of Sciences, Engineering, and Medicine and others. 2018. *Assessing and responding to the growth of computer science undergraduate enrollments*. National Academies Press.
- [26] Nhan Nguyen and Sarah Nadi. 2022. An empirical evaluation of GitHub copilot's code suggestions. In *Proceedings of the 19th International Conference on Mining Software Repositories*. ACM, Pittsburgh Pennsylvania, 1–5. <https://doi.org/10.1145/3524842.3528470>
- [27] Clíodhna O'Connor and Helene Joffe. 2020. Intercoder reliability in qualitative research: debates and practical guidelines. *International journal of qualitative methods* 19 (2020), 1609406919899220.
- [28] Zachary A. Pardos and Shreya Bhandari. 2023. Learning gain differences between ChatGPT and human tutor generated algebra hints. arXiv:2302.06871 [cs.CY]
- [29] James Prather, Paul Denny, Juho Leinonen, Brett A. Becker, Ibrahim Albluwi, Michelle Craig, Hieke Keuning, Natalie Kiesler, Tobias Kohn, Andrew Luxton-Reilly, Stephen MacNeil, Andrew Peterson, Raymond Pettit, Brent N. Reeves, and Jaromir Savelka. 2023. The Robots are Here: Navigating the Generative AI Revolution in Computing Education. arXiv:2310.00658 [cs.CY]
- [30] James Prather, Brent N. Reeves, Paul Denny, Brett A. Becker, Juho Leinonen, Andrew Luxton-Reilly, Garrett Powell, James Finnie-Ansley, and Eddie Antonio Santos. 2023. "It's Weird That it Knows What I Want": Usability and Interactions with Copilot for Novice Programmers. <http://arxiv.org/abs/2304.02491> arXiv:2304.02491 [cs].
- [31] Sami Sarsa, Paul Denny, Arto Hellas, and Juho Leinonen. 2022. Automatic Generation of Programming Exercises and Code Explanations Using Large Language Models. In *Proceedings of the 2022 ACM Conference on International Computing Education Research V.1*. ACM, Lugano and Virtual Event Switzerland, 27–43. <https://doi.org/10.1145/3501385.3543957>
- [32] Jaromir Savelka, Arav Agarwal, Marshall An, Chris Bogart, and Majd Sakr. 2023. Thrilled by Your Progress! Large Language Models (GPT-4) No Longer Struggle to Pass Assessments in Higher Education Programming Courses. In *Proceedings of the 2023 ACM Conference on International Computing Education Research V.1* (ICER 2023). ACM. <https://doi.org/10.1145/3568813.3600142>
- [33] Jaromir Savelka, Arav Agarwal, Christopher Bogart, and Majd Sakr. 2023. Large Language Models (GPT) Struggle to Answer Multiple-Choice Questions about Code. arXiv:2303.08033 [cs.CL]
- [34] Jaromir Savelka, Paul Denny, Mark Liffiton, and Brad Sheese. 2023. Efficient Classification of Student Help Requests in Programming Courses Using Large Language Models. arXiv:2310.20105 [cs.CY]
- [35] Haoye Tian, Weiqi Lu, Tsz On Li, Xunzhu Tang, Shing-Chi Cheung, Jacques Klein, and Tegawendé F. Bissyandé. 2023. Is ChatGPT the Ultimate Programming Assistant – How far is it? <http://arxiv.org/abs/2304.11938> arXiv:2304.11938 [cs].